

Chapter 6

Binary Search Trees and Linear Time Sorts.

6.1 Linear Time Sorts

Counting sort. Counting sort assumes that each of the n input elements is an integer at the size at most k . It runs in $\Theta(n + k)$ time, so that when $k = O(n)$, counting sort runs in $\Theta(n)$ time. Counting sort first determines, for each input element x , the number of elements less than or equal to x . It then uses this information to place element x directly into its position in the output array. For example, if 17 elements are less than or equal to x , then x belongs in output position 17. We must modify this scheme slightly to handle the situation in which several elements have the **same value**, since we do not want them all to end up in the same position.

```
1 let B and C be new arrays at size  $n$  and  $k$ 
2 for  $i \in [0, k]$  do
3   |  $C_i \leftarrow 0$ 
4 end
5 for  $j \leftarrow [1, n]$  do
6   |  $C_{A_j} \leftarrow C_{A_j} + 1$ 
7 end
8 for  $i \in [1, k]$  do
9   |  $C_i \leftarrow C_i + C_{i-1}$ 
10 end
11 for  $i \in [n, 1]$  do
12   |  $B_{C_{A_j}} \leftarrow A_j$ 
13   |  $C_{A_j} \leftarrow C_{A_j} - 1$  // to handle duplicate values
14 end
15 return  $B$ 
```

Notice that the Counting sort can beat the lower bound of $\Omega(n \log n)$ only because it is not a comparison sort. In fact, no comparisons between input elements occur anywhere in the code. Instead, counting sort uses the actual values of the elements to index into an array.

An important property of the counting sort is that it is **stable**.

Stable Sort.

We will say that a sorting algorithm is stable if elements with the same value appear in the output array in the same order as they do in the input array.

Counting sort's stability is important for another reason: counting sort is often used as a subroutine in radix sort. As we shall see in the next section, in order for radix sort to work correctly, counting sort must be stable.

Radix sort Radix sort is the algorithm used by the card-sorting machines you now find only in computer museums. The cards have 80 columns, and in each column, a machine can punch a hole in one of 12 places. The sorter can be mechanically "programmed" to examine a given column of each card in a deck and distribute the card into one of 12 bins depending on which place has been punched. An operator can then gather the cards bin by bin, so that cards with the first place punched are on top of cards with the second place punched, and so on.

The Radix-sort procedure assumes that each element in the array A has d digits, where digit 1 is the lowest-order digit and digit d is the highest-order digit.

```

1 for  $i \in [1, d]$  do
2   | use a stable sort to sort array  $A$  on digit  $i$ 
3 end
```

Correctness Proof. By induction on the column being sorted.

- Base. Where $d = 1$, the correctness follows immediately from the correctness of our base sort subroutine.
- Induction Assumption. Assume that Radix-sort is correct for any array of numbers containing at most $d - 1$ digits.
- Step. Let A' be the algorithm output. Consider $x, y \in A$. Assume without losing generality that $x > y$. Denote by x_d, y_d their d -digit and by $x_{/d}, y_{/d}$ the numbers obtained by taking only the first $d - 1$ digits of x, y . Separate in two cases:
 - If $x_d > y_d$ then a scenario in which x appear prior to y is imply contradiction to the correctness of our subroutine.
 - So consider the case in which $x_d = y_d$. In that case, it must hold that $x_{/d} > y_{/d}$. Then the appearance of x prior to y either contradicts the assumption that the base algorithm we have used is stable or that x appears before y at the end of the $d - 1$ iteration. Which contradicts the induction assumption.

The analysis of the running time depends on the stable sort used as the intermediate sorting algorithm. When each digit lies in the range 0 to $k-1$ (so that it can take on k possible values), and k is not too large, counting sort is the obvious choice. Each pass over n d -digit numbers then takes $\Theta(n + k)$ time. There are d passes, and so the total time for radix sort is $\Theta(d(n + k))$.

Bucket sort. Bucket sort divides the interval $[0, 1)$ into n equal-sized subintervals, or buckets, and then distributes the n input numbers into the buckets. Since the inputs are uniformly and independently distributed over $[0, 1)$, we do not expect many numbers to fall into each bucket. To produce the output, we simply sort the numbers in each bucket and then go through the buckets in order, listing the elements in each.

```

1 let B[0 : n - 1] be a new array for  $i \leftarrow [0, n-1]$  do
2   | make  $B_i$  an empty list
3 end
4 for  $i \leftarrow [1, n]$  do
5   | insert  $A_i$  into list  $B_{\lfloor nA_i \rfloor}$ 
6 end
7 for  $i \leftarrow [0, n-1]$  do
8   | sort list  $B_i$ 
9 end
10 concatenate the lists  $B_0, B_1, \dots, B_{n-1}$  together and
11 return the concatenated lists

```

6.2 Binaries Trees and How to Encode Them.

We have already seen, in heaps, that organizing our data in a graph-like structure can offer a speed advantage. For future applications, and in particular for maintaining data in sorted order, we will have to encode our data using binary trees. These trees may not be almost complete and also have to support pointer manipulations, specifically placing a binary tree as a left or right subtree of a given node. To enable this, we will have to treat the **right**, **left**, and **parent** as variables, in contrast to heaps where they are determined completely by the node index. We begin this section by stating definitions.

Definition 6.2.1. 1. *Binary Tree:* A tree in which any vertex has at most two children.

2. A descendant of vertex x is a vertex in the subtree whose root is x . A left descendant of vertex x is either a vertex in the subtree whose root is the left child of x , or x 's left child.

3. An ancestor of x is a vertex to which x belongs as a descendant.

4. A leaf is a vertex without children.

5. Height of vertex x is the length of the longest simple path (without cycles) between x and one of the leaves.
6. Height of the tree is the height of its root, which is usually denoted by h .

We encode a binary tree by associating a field to each vertex x , representing its right, left children, and parent. We use the notation $x.\text{left}$ to refer to the left child of x , although the physical implementation may differ conceptually. For example, the way binary trees are implemented in Cormen is through 4 arrays. The first stores the value of x , while the others store pointers of specific types. For instance, the array LEFT, where $\text{LEFT}.x$ stores the left child of x .

If nothing else has been mentioned, then we can assume that we can add additional fields to the vertices.

6.3 Binary Search Trees.

Binary search tree (BST) is a binary tree which any node x of it:

1. Contains a field key, storing a number $x.\text{key}$.
2. Any left descendant y of x satisfies $y.\text{key} \leq x.\text{key}$.
3. Any right descendant y of x satisfies $y.\text{key} \geq x.\text{key}$.

Question. Let T be a binary search tree, Where are the minimum and maximum values of T ? (most left and right nodes).

Definition 6.3.1. Let T be a binary search tree, and let x be a node belonging to it. The predecessor of x will be defined as a vertex y such that $y.\text{key} \leq x.\text{key}$ and $y.\text{key}$ is maximal among the nodes satisfying this condition. If we were to set the values of T in sorted order, then the predecessor of x would be located on its left. The successor of x will be defined as y , where x is the predecessor of y .

Functionality of BST.

1. Search(T , key): returns a pointer to the vertex whose key equals key.
2. Min(T): returns a pointer to the vertex with the minimum value in T .
3. Max(T): returns a pointer to the vertex with the maximum value in T .
4. Predecessor(x): returns a pointer to the predecessor of x .
5. Successor(x): returns a pointer to the successor of x .
6. Insert(T , key): inserts key into T (creates a new vertex).
7. Delete(T , x): removes x from T .
8. Inorder(T): outputs T 's keys in sorted order.

Data: T - tree, key - key to search for

Result: pointer to vertex with key equal to key

```

1 Function Search( $T, key$ )
2   if  $T$  is empty then
3     |   return null;
4   end
5   if  $T.key$  equals  $key$  then
6     |   return  $T$ ;
7   end
8   if  $key$  is less than  $T.key$  then
9     |   return Search( $T.left, key$ );
10  end
11  else
12    |   return Search( $T.right, key$ );
13  end
14 end

```

Data: T - tree

Result: pointer to vertex with minimum value in T

```

1 Function Min( $T$ )
2   if  $T$  is empty then
3     |   return null;
4   end
5   if  $T.left$  is empty then
6     |   return  $T$ ;
7   end
8   return Min( $T.left$ );
9 end

```

Data: T - tree

Result: pointer to vertex with maximum value in T

```

1 Function Max( $T$ )
2   if  $T$  is empty then
3     |   return null;
4   end
5   if  $T.right$  is empty then
6     |   return  $T$ ;
7   end
8   return Max( $T.right$ );
9 end

```

Data: x - vertex

Result: pointer to predecessor of x

```

1 Function Predecessor( $x$ )
2   if  $x.left$  is not empty then
3     |   return Max( $x.left$ );
4   end
5    $y \leftarrow x.parent$ ;
6   while  $y$  is not empty and  $x$  is  $y.left$  do
7     |    $x \leftarrow y$ ;
8     |    $y \leftarrow y.parent$ ;
9   end
10  return  $y$ ;
11 end

```

Data: x - vertex

Result: pointer to successor of x

```

1 Function Successor( $x$ )
2   if  $x.right$  is not empty then
3     |   return Min( $x.right$ );
4   end
5    $y \leftarrow x.parent$ ;
6   while  $y$  is not empty and  $x$  is  $y.right$  do
7     |    $x \leftarrow y$ ;
8     |    $y \leftarrow y.parent$ ;
9   end
10  return  $y$ ;
11 end

```

Data: T - tree, key - key to insert

Result: inserts key into T

```

1 Function Insert( $T, key$ )
2    $newNode \leftarrow$  create new vertex with key  $key$ ;
3    $y \leftarrow$  null;
4    $x \leftarrow T$ ;
5   while  $x$  is not empty do
6      $y \leftarrow x$ ;
7     if  $key$  is less than  $x.key$  then
8        $x \leftarrow x.left$ ;
9     end
10    else
11       $x \leftarrow x.right$ ;
12    end
13  end
14   $newNode.parent \leftarrow y$ ;
15  if  $y$  is null then
16     $T \leftarrow newNode$ ;
17  end
18  else if  $key$  is less than  $y.key$  then
19     $y.left \leftarrow newNode$ ;
20  end
21  else
22     $y.right \leftarrow newNode$ ;
23  end
24 end

```

Data: T : The input tree

```

1 if  $T$  is not empty then
2   Inorder( $T.left$ );
3   Output  $T.key$ ;
4   Inorder( $T.right$ );
5 end

```

Algorithm 1: Inorder(T)