



Hashing

Up to this point, all the data structures we have seen in this course assume nothing but comparability about the input keys they are expected to manage. Hash functions, in general, are functions¹ from the key space into a lower dimensional space, which can be thought of as our memory storage.

Example 1. For example, assume the keys are taken from the integers, and the hash function $h : \mathbb{N} \rightarrow [10]$ sends numbers to their residue modulo 10, namely $h(x) = x \bmod 10$. Now, one might use h and a 10-length array T to store numbers. Any time he would like to insert a number x into the structure, he would set $T[h(x)] \leftarrow x$.

However, this is not a “good” method, since two different keys with the same residue modulo 10 are mapped into the same cell in T , which we count as a collision:

Definition 1 (collision.). Given a function $h : U \rightarrow \star$, we name any pair of different keys $x \neq y \in U$ that are mapped by h to the same value, namely $h(x) = h(y)$, a collision.

Clearly, if we assume that no collisions are going to occur, i.e. all the given keys through the running of the program $x_1, x_2, \dots, x_n$ satisfy that for any $x_i \neq x_j$, it follows that $h(x_i) \neq h(x_j)$, then the method in Example 1 provides a data structure that² supports access, insertion, and deletion in constant time. This gives us the intuition that the complexity of a general data structure which uses a hash function depends on the way it resolves collisions. In this recitation, we will present ways to handle collisions. The first is a heuristic called **open addressing** (or **closed hashing**). The second, **open hashing**, will also be analyzed, and we will show a characterization that, if satisfied, then we get a good running time in expectation. Finally, we will also show examples of function families that satisfy the characterization and examples of families that do not.

1 Open Addressing (Closed Hashing)

Open addressing uses a static m -length array T , hash function h , and handles collisions by probing cells in a specific order. For example, linear probing will first check if the cell at position $h(x)$ is not occupied or already contains x . If it is, then $h(x)$ is considered the correct place to store x . Otherwise, the next cell is probed, namely the cell at position $h(x) + 1$, and so on.

¹projections.

²Again, relative to assumptions on the input.

```

1 let  $i \leftarrow 0$ 
2 while  $T[h(x) + i] \neq \text{Null}$  do
3   |  $i \leftarrow i + 1$ 
4 end
5 return  $h(x) + i$ 

```

Algorithm 1: linear probing - access

2 Universal Hashing

Now that it is understood that collisions govern the running time, we are ready to study a classification of mappings which generate a low number of collisions in expectation.

Definition 2. Let $\mathcal{H} = \{h_i : U \rightarrow [m]\}$ be a family of function from the domain U into $[m] = \{0, 1, \dots, m-1\}$. $\mathcal{H}$ will be said universal if for any $x \neq y \in U$:

$$\Pr_{h \sim \mathcal{H}} [h(x) = h(y)] \leq \frac{1}{m}$$

Question. For $x = y$ what is the probability that $h(x) = h(y)$?

Example 2. $\mathcal{H}$ is the set of all function from $U \rightarrow [m]$.

Picking a function randomly from the set of all functions is equivalent to picking the targets $h(x)$, for each x , independently. So the probability that $h(x) = h(y)$ is as exactly as the probability, in the balls to bins experiment, that two specific balls (x and y in our case) will be thrown into the same bin. As we computed in the last recitation, that probability equals $1/m$.

Example 3. $\mathcal{H}$ is all the inner products by a n -length binary vector. Namely,

$$\mathcal{H} = \{h_a : \mathbb{F}_2^n \rightarrow \mathbb{F}_2, h_a(x) = \langle a, x \rangle \mid a \in \mathbb{F}_2^n\}$$

Observe that in $\mathbb{F}_2$, the equivalence $h(x) = h(y) \Leftrightarrow h(x) + h(y) = 0$ holds. If $x \neq y$, then there is at least one coordinate i such that $x_i \neq y_i$. In other words, $x_i + y_i = 1$. Thus:

$$\begin{aligned} h_a(x) + h_a(y) &= \langle a, x + y \rangle = a_i(x_i + y_i) + \sum_{j \neq i} a_j(x_j + y_j) \\ &= a_i + \sum_{j \neq i} a_j(x_j + y_j) \end{aligned}$$

Thus:

$$\begin{aligned}
\Pr[h_a(x) = h_a(y)] &= \sum_{z \in \mathbb{F}_2} \Pr \left[a_i = z \cap \left\{ \sum_{j \neq i} a_j(x_j + y_j) = z \right\} \right] \\
&= \sum_{z \in \mathbb{F}_2} \Pr[a_i = z] \cdot \Pr \left[\left\{ \sum_{j \neq i} a_j(x_j + y_j) = z \right\} \right] \\
&= \sum_{z \in \mathbb{F}_2} \frac{1}{2} \cdot \Pr \left[\left\{ \sum_{j \neq i} a_j(x_j + y_j) = z \right\} \right] \\
&= \frac{1}{2} \sum_{z \in \mathbb{F}_2} \Pr \left[\left\{ \sum_{j \neq i} a_j(x_j + y_j) = z \right\} \right] = \frac{1}{2}
\end{aligned}$$

When in the second pass we used independence, and in the last, we used that probabilities sum to 1.

Example 4. $\mathcal{H}$ is the set of all binary matrices : $\mathbb{F}_2^n \rightarrow \mathbb{F}_2^k$.

Observe that the equivalence $h(x) = h(y)$ implies an equivalence in any coordinate. Since the i th coordinate of $h(x)$, namely $h(x)_i$, is given by an inner product of the i th row of h , which is picked randomly from the n -length binary vectors, we get by Example 3 that the probability of equivalence in the i th coordinate is $\frac{1}{2}$, and since the rows are picked independently, we get:

$$\Pr[h(x) = h(y)] = \Pr \left[\bigcap_i^k h(x)_i = h(y)_i \right] = \left(\frac{1}{2} \right)^k$$

Exercise 1. Let $\mathcal{H}$ be a universal family from $U \rightarrow [m]$. Show that $\mathcal{H}^* : U \times U \rightarrow [m]$ is defined by:

$$h^*(x, y) = h(x) + h(y) \text{ for } h \in \mathcal{H}$$

is not universal.

Solution 1. Let $x, y \in U$. Consider the elements $z_1 = (x, y)$ and $z_2 = (y, x)$. On one hand, $z_1 \neq z_2 \in U \times U$, on the other hand, for any $h^* \in \mathcal{H}^*$ (defined with some $h \in \mathcal{H}$) it holds that:

$$h^*(z_1) = h^*((x, y)) = h(x) + h(y) = h(y) + h(x) = h^*((y, x)) = h^*(z_2)$$

So, we get that with probability 1, $h(z_1) = h(z_2)$, and hence $\mathcal{H}^*$ is not universal.

Equipped with universal hash functions, we can now say intelligent 'things' about the expected number of collisions. Specifically, denote by $h \sim \mathcal{H}$ a sampled hash function. Let us denote by I the input: $I = \{x_1, x_2, \dots, x_n\} \subset U$. We can define for each $x \in I$ the random variable C_x , which counts the number of collisions in I involving x . Namely:

$$C_x = |\{y \in I : h(x) = h(y)\}|$$

Note, How to recall: at least 1 element for x itself, and $\frac{n-1}{m}$ is what we would guess to have by throwing $n - 1$ balls to m bins.

Claim 1. If $\mathcal{H}$ is universal hash family, then $\mathbf{E}[C_x] = 1 + \frac{n-1}{m}$.

Proof. For any $x, y \in I$, let $X_{x,y}$ be the indicator of the event $h(x) = h(y)$. Observe that $C_x = \sum_{y \in I} X_{x,y}$. Therefore, by the linearity of expectation and $\mathcal{H}$ being a universal hash family, we have:

$$\begin{aligned} \mathbf{E}[C_x] &= \mathbf{E}\left[\sum_{y \in I} X_{x,y}\right] = \sum_{y \in I} \mathbf{E}[X_{x,y}] = \sum_{y=x \in I} \mathbf{E}[X_{x,y}] + \sum_{y \neq x \in I} \mathbf{E}[X_{x,y}] \\ &= 1 + \sum_{y \neq x \in I} \Pr[h(x) = h(y)] \leq 1 + \frac{n-1}{m} \end{aligned}$$

□

2.1 Universal Hashing Table (Open Hashing)

Universal Hashing Table also defined by an m -array (table), and an hash function h which index keys $x \in U$ to $T[h(x)]$ cells, Yet in contrast to open addressing each cell might contains multiple keys by maintaining a linked list. Accessing to a stored key, or checking if it is in T , is done by iterating the linked list at suited cell. Appending a key, is done by first checking the key is not already stored in the table and then appending it to the tail of the suited linked list. Line 2 exhibits the access routine, you are encouraged to complete the others.

Question. Fix a table, and consider $x \in U$. What will be the running time of each of the operation: Insert x , Access x , Delete x ? (Linear at the number of keys in the list as $T[h(x)]$).

Question. In hashing terms, what is the length of $T[h(x)]$ equal to? The number of collision with x .

Hence, we infer that the expected running time is bounded by the expected number of collision with x . If h is sampled uniformly at random from a universal

```

1 linked ← T[h(x)]
2 for node ∈ linked do
3   | if node.value = x then
4   |   | return node
5   | end
6 end
7 return last node
```

Algorithm 2: Universal Hashing Table - access

family $\mathcal{H}$, then by Claim 1, the expected running time is $1 + \frac{n-1}{m} = O(1)$.

2.2 How Hashing is Implemented in Python?

In most modern programming languages, the Dictionary data structure is implemented by a variation of Cuckoo hashing, which works as follows: We sample

two functions from $\mathcal{H}$, denoted by h_1 and h_2 . Checking if an element x is in the table is done by checking if it is in either $T[h_1(x)]$ or $T[h_2(x)]$. Inserting it into the table is done by Line 3.³

```

1 if either  $T[h_1(x)]$  or  $T[h_2(x)]$  are empty then
2   |   Pick an empty spot and set  $x$  there.
3   |   return
4 end
5 else
6   |    $i \leftarrow$  flip a coin.
7   |    $z \leftarrow h_i(x)$ 
8   |   set  $T[h_i(x)] \leftarrow x$ 
9   |   Call Insert( $z$ ).
10 end
```

Algorithm 3: Cuckoo - Insert

Roughly speaking, the algorithm fails when we close a cycle, namely after several recursive calls, back to the point where we were called to insert x . The idea of the analysis is that if the functions h_1, h_2 were picked randomly, or very similar to randomly (i.e., from k -universal families), that event is unlikely, and the number of scanned spots is constant.

³Here we lie, a slightly clever (but not too mach) rule is used to pick i .